

Shivam Pandey

+91-8401260284 | shivam574488@gmail.com | [linkedin](#) | [github](#) | [Huggingface](#)

SUMMARY

AI Engineer with 2+ years of experience building production RAG systems, AI agents, and LLM pipelines at Jellyfish Technologies. Skilled in agentic workflows, retrieval-augmented generation, and deploying AI across customer support, voice interfaces, and document processing.

EXPERIENCE

AI Engineer

Feb 2024 – Present

Jellyfish Technologies

Noida, India

- Built model evaluation pipelines to benchmark LLM performance across domain-specific tasks.
- Designed and deployed a **Freshdesk AI Agent with RAG** using Vertex AI, handling **70–100 tickets daily** with **45% auto-resolution rate**, reducing manual workload significantly.
- Optimized the retrieval system by compressing and structuring Freshdesk solutions and forums, cutting storage costs while improving search relevance.
- Developed a **real-time analytics dashboard** tracking resolution rates, turnaround times, and auto-response performance.
- Built an end-to-end SFT data pipeline for a legal domain LLM — scraped and parsed Supreme Court and High Court judgements from 1800–2026, cleaned raw judicial text, and transformed case data into instruction-response pairs for supervised fine-tuning.

Machine Learning Research Assistant

May 2023 – July 2024

KCC ITM

Greater Noida, UP

- Explored deep learning techniques to generate maze puzzles using GANs and VAEs to create solvable mazes with diverse patterns.
- Developed an interactive game in Python to test the generated mazes.

EDUCATION

Abdul Kalam Technical University

Greater Noida, UP

Bachelor of Technology in Computer Science and Engineering

Nov. 2020 – July 2024

PROJECTS

Emotion Classifier DistilBERT Fine-Tuning | *PyTorch, Hugging Face Transformers* | Huggingface

Jan 2023

- Fine-tuned DistilBERT for **emotion classification** on transcripts and reviews, enabling tone-aware auditing. Delivered a scalable pipeline and open-source model hosted on Hugging Face Hub.

1N20 Home Services — AI Voice Sales Assistant | *OpenAI Agents SDK, Pinecone, FastAPI, MongoDB*

Jan 2026 – Present

- Built a dual-agent system with handoff, 12+ tools, shared memory, and session continuity for voice-based customer sales.
- Developed a semantic search pipeline using Pinecone to retrieve products from **9,000+ indexed items** with image-augmented responses.
- Built a real-time voice interface that captures leads, stores them in MongoDB, and triggers automated follow-up emails.

Levett Consultancy — IT Support Automation | *Vertex AI, LangChain, Google Cloud, Guardrails AI*

Jul 2025 – Present

- Built an email triage system that classifies ticket complexity and routes queries — achieving **94% routing accuracy**. Developed an AI response system using Vertex AI that searches a 15-year knowledge base to answer L1 queries with **87% accuracy**.
- Added a Guardrails AI layer to ensure all auto-responses are grounded in retrieved content before delivery.

Document Intelligence System | *Neo4j, LLMs, AWS Textract, Qdrant*

Jan 2025

- Built a document processing system that maps extracted data into a Neo4j knowledge graph, reducing manual work by 50%. Designed an OCR pipeline using AWS Textract to extract structured data from scanned files for graph-based entity analysis.

TECHNICAL SKILLS

Languages: Python, C/C++, SQL, JavaScript/TypeScript

Database: PostgreSQL, MongoDB, Neo4j, MySQL, Redis, Milvus, Pinecone

Libraries/Frameworks: NumPy, Plotly, Pandas, Scikit-Learn, PyTorch, Streamlit, FastAPI, Flask, pytest, Node.js, React.js, Jest, Hugging Face Transformers, MLflow

Gen AI: Prompt Engineering, AutoGen, LangGraph, LangChain, LangSmith, LlamaIndex, CrewAI, Unsloth, N8n

Developer Tools: Git, GitHub, VertexAI, Azure AI, Docker, AWS

Certifications: Databricks Certified ML Professional, Foundation of Generative AI, Data Engineering (IBM), Tools for Data Science